


# Lasso predictions scale better with noise

Prediction error vs number of noise parameters



Note: This summarizes the results of a logistic regression with and without penalization for different number of predictors. The x-axis has the total number of predictors, only 10 of which are predictive. The y-axis is the RMSE between the fit expectations of  $y_i$  and the true expectation of  $y_i$